

# 1 Sparse Matrix Multiplication

The algorithmic task in this question is to multiply two  $n \times n$  sparse matrices  $A$  and  $B$ , to obtain a result  $C \leftarrow AB$ . By sparse, we mean that almost all of the entries in  $A$  and  $B$  are 0.

Each matrix is given as a list of triples  $(i, j, x)$ . We have  $A[i][j] = x$  if  $(i, j, x)$  is in the list, and  $A[i][j] = 0$  otherwise.

Assume that we index beginning in the top left corner, and the matrices are 0-indexed (for example, this means that the top left corner of  $A$  is  $A[0,0]$ , and the top right corner of  $A$  is  $A[0, n-1]$ ). The triples can, of course, be given in any order. However, we assume that the entries of  $A$  are listed in column-major order: entries are listed one column after the other, with entries within a column listed row-by-row. Similarly, the entries of  $B$  are listed in row-major order: one row at a time; entries within a row sorted by column.

For example, assume we are computing the product

$$\begin{matrix} \begin{bmatrix} 0 & 0 & 4 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 10 & 0 \\ 0 & 4 & 0 & 1 \end{bmatrix} & \begin{bmatrix} 0 & 0 & 0 & 0 \\ 9 & 0 & 0 & 0 \\ 0 & 7 & 6 & 0 \\ 0 & 0 & 3 & 0 \end{bmatrix} & = & \begin{bmatrix} 0 & 28 & 24 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 70 & 60 & 0 \\ 36 & 0 & 3 & 0 \end{bmatrix} \\ A & B & & C \end{matrix}$$

$A$  would be represented using the triples  $(3, 1, 4), (0, 2, 4), (2, 2, 10), (3, 3, 1)$  and  $B$  would be represented as  $(1, 0, 9), (2, 1, 7), (2, 2, 6), (3, 2, 3)$ . The solution to this instance would be

$$C = (3, 0, 36), (0, 1, 28), (0, 2, 24), (2, 1, 70), (2, 2, 60), (3, 2, 3).$$

Consider an algorithm that goes through each nonempty column of  $A$  and each nonempty row of  $B$ , performing all relevant multiplications and storing the result in  $C$ . We assume that  $A$  and  $B$  are both  $n \times n$  matrices. Recall that matrix multiplication can be defined as  $C[i][j] = \sum_k A[i][k] \cdot B[k][j]$ ; we use this notation in our pseudocode for clarity.

Note that this algorithm leaves  $C$  as a list of possibly-duplicate tuples. A full matrix-multiplication algorithm would further process  $C$  to remove entries that correspond to the same cell (i.e. if  $(1, 1, 3)$  and  $(1, 1, 5)$  were both in  $C$ , we would want to add them to obtain  $(1, 1, 8)$ ). In Algorithm 1 we ignore this step for simplicity.

## Implement

---

### Algorithm 1 Sparse Matrix Multiplication

---

```

for  $k = 1$  to  $n$  do                                     ▷ Column  $k$  of  $A$ ; row  $k$  of  $B$ 
  for  $i = 1$  to  $n$  do ▷ These loops iterate over every pair of nonzero items in this row/column
    for  $j = 1$  to  $n$  do
       $c \leftarrow A[i][k] \cdot B[k][j]$ 
      Append  $(i, j, c)$  to  $C$ 

```

---

1. Implement the above sparse matrix multiplication algorithm efficiently.

By efficient, we mean that your running time should be proportional to the number of non-zero entries in the output (elementary products to be summed to entries in  $C$ ); in particular it should not be  $\Theta(n^3)$  for sufficiently sparse matrices.

$A$  and  $B$  are given as lists of triples, one triple per line. The program should take five arguments:  $n$  (the size of the matrices), the number of entries in  $A$ , the name of the file containing  $A$ , the number of entries in  $B$ , and finally the name of the file containing  $B$ . The first CodeJudge test, which corresponds to the above example, has arguments

```
4 4 A.txt 4 B.txt
```

We have given exercises on CodeJudge to allow you to verify your implementation.

### Questions to Discuss

2. Give an example where  $C$  *does not* need a final summation step (all entries are for distinct positions in  $C$ ). Give an example where  $C$  *does* need a final summation step.
3. Design and carry out an experiment to compare the performance of your sparse matrix multiplication code with a dense matrix multiplication implementation.